
Automatic Recognition of Pictured Dishes in Food-101

Project Update - 1

Abishek Kumar

Department of Computer Sciences
University of Wisconsin-Madison
akumar225@wisc.edu

1 Introduction

Classification is an important tool in today's world where big data is used to make all kinds of decisions in economics, medicine, and more. For example, detecting spam in emails can be identified as a classification problem with two classes as spam and not spam. The problem of object detection and classification is a widely researched topic in machine learning due to its broad range of applications. There is an array of different algorithms that have been developed over the years to solve this problem. Food classification in particular has many applications like helping patients track their calorie intake, auto-organizing of pictures in mobile phones, etc. But, unlike other image recognition tasks wherein we have definitive features separating each class, the problem of food classification is unique due to its high intra-class variability. Lighting, orientation, and the very realization of a recipe are some of the factors which can contribute to this variability.

In this project, we will try to understand the performance of simple classification algorithms like Linear Regression, KNN for the food classification task and compare their performance against more sophisticated algorithms like Neural Networks. We will use the Food-101 dataset for training our classifiers and evaluate their performance using holdout method, precision/recall, and ROC curve. Finally, we will propose a recommended algorithm based on our results.

2 Dataset: Food - 101

The dataset consists of 101 food categories with 101,000 images, 1000 images for each category split into 750 training images and 250 test images. Both the training and test images have been downloaded from foodspotting.com. The dataset has the top 101 food categories from the website. The test images have been manually cleaned, while the training images were left intentionally unclean and have some amount of noise. This noise comes mostly in the form of intense colors and sometimes mislabelled samples. All images have been re-scaled to have a maximum side length of 512 pixels and smaller images have been discarded.

2.1 Preprocessing

As a pre-processing step, we first separated the training and testing images into two different folders. Then, we resized each image in the dataset to the size of 128 x 128 x 3. We chose this size mainly due to memory constraints. Initially, we tried to experiment with the size of 299 x 299 x 3. But the number of pixels in this case was huge - 268,203. This meant that the model had to learn 268,203 weights for fitting the data. This led to memory issues even on a machine with 500GB of RAM. Hence, we decided to resize the images to 128 x 128 x 3 which gave a fair trade-off between image quality and feasibility of computation. We used python's cv2 library for resizing the images.

Table 1: Training and Testing Accuracy for each classifier

| Sno | Loss Function | Lambda | Training Accuracy | Testing Accuracy |
|-----|---------------|--------|-------------------|------------------|
| 1 | Squared-Loss | 0.001 | 1.08% | 1.36% |
| 2 | Log Loss | 0.001 | 4.8% | 1.4% |
| 3 | Hinge | 0.001 | 3.8% | 1.5% |
| 4 | Hinge | 10 | 4.1% | 1.9% |
| 5 | Hinge | 30 | 3.2% | 1.5% |
| 6 | Hinge | 10 | 4.8% | 3.3 % |

3 Experimental Setup

The experiments were done on ml.p3.16xlarge machine on Amazon AWS Sagemaker. The machine has the following configuration - 64 cores, 500GB of RAM and 250GB of hard disk space. We started with the following machines at first - 16GB RAM, 4 cores and 31GB RAM, 8 cores. But, both these machines were not able to handle the huge size of the dataset and hence we decided to go for a bigger machine. The entire implementation was done in Python and we used Python's scikit-learn library for training and evaluating the models.

4 Implementation

For the first phase of the project, we perform the classification using ridge regression. Sklearn library provides a direct implementation for this - RidgeClassifier. This implementation loads the entire dataset into memory and trains the model. We could not use it directly for this dataset as it was running into memory issues. Hence, we decided to train the model iteratively using Sklearn's Stochastic Gradient Descent Classifier (SGDClassifier). We split the dataset into batches of size 750 and iteratively trained the model. For testing purpose, we split the data into batches of size 250 and computed the accuracy in each batch. Finally, we averaged the accuracy across all batches to produce overall testing accuracy.

The entire implementation of the project can be found in this Github Repository - [Automatic-Recognition-of-Pictured-Dishes-in-Food-101](#).

5 Results

We trained 5 different classifiers with different settings for the loss function and the regularization parameter lambda. These classifiers were trained on a subset of data - 7500 training images and 2500 testing images to quickly identify the optimal setting for the loss function and the regularization parameter. The training and testing images were randomly sampled from the original dataset. Once we identified the optimal setting, we also trained a 6th classifier on the entire dataset using the optimal setting. Table 1 gives the training and testing accuracy for each classifier.

6 Conclusion

From Table 1, we can see that both the training and testing accuracy are poor for all 6 classifiers. This could be due to the fact that the classifier considers each pixel as an individual feature and ignores the spatial correlation between pixels. Also, the number of training samples for each class (750) is very less compared to the number of features (49,152). This is an underdetermined system and hence the model is not able to learn the weights accurately. We are also evaluating accuracy score based on the top-1 prediction. This is a very harsh metric as it expects for each sample that the correct label is predicted. Instead, we can also get the score for the class being in top-3 / top-5 predictions and evaluate accordingly.

7 Next Steps

- November 29th: Evaluate the top-3 and top-5 prediction accuracy for the RidgeClassifiers. As per the initial project proposal, apply KNN Algorithm for classification and produce results.
- December 7th: Develop a neural network for classification and produce results.
- December 14th: Compare the results between all three algorithms and prepare final report.